

Machine Learning: Fundamentals and Applications

House Price Prediction

Group Members

Harsh Agrawal
Jayraj Rameshbhai Kanani
Nidhi Singh
Sophia Roper





TABLE OF CONTENTS

01

DATASET

Discussion of dataset
columns and attributes

02

Problem Statement

Why is there a need for the
analysis / model

03

EDA

Exploratory data analysis to
find useful patterns

04

Heatmap

Find correlation and useful
attributes

05

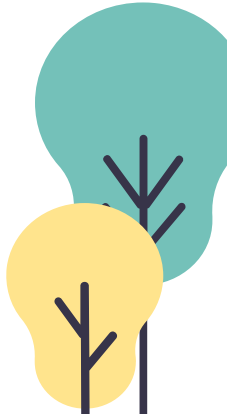
Regression Training

Using different regression
algorithms to learn the data

06


Conclusion

Conclusions drawn from the
data and the models built



Dataset

- This dataset comprises 79 explanatory variables that cover practically every facet of Ames, Iowa's residential dwellings.
- Has 1460 rows and 79 columns related to residential homes
- Target value is Sales Price to predict final price of each home



House price
prediction


Problem Statement

When you ask a home buyer to describe their ideal home, they are unlikely to start with the basement ceiling height or the closeness to an east-west railroad. However, the data from this playground competition shows that far more factors influence price negotiations than the number of bedrooms or the presence of a white-picket fence. However, the data in this thesis shows that the number of bedrooms and floors have a greater impact on the price of a home. I also want to use this dataset to anticipate an acceptable home price based on these attributes of the properties.



GOALS

The goal of our Project is to train machine learning models to predict the housing prices and find which aspects of the house influence the housing prices mostly.



Correlation Matrix

Obtained correlation between two variable to check the highly correlated variable with its dependent variable

Dropping Columns

Columns with mostly NaN values were dropped, including FireplaceQu, Alley, Id, PoolQC, Fence, and MiscFeature


Using only necessary columns

Eliminated all the least important columns from the dataset



Scatter plot of the sale price

Most of the points are assembled on the bottom. And there seems to be no large outliers in the sale price variable



Heatmap

Observation

- No direct patterns are observed from heatmap.
- Continuous numeric values seem to affect correlation most.
- The column “OverallQual” seems to be directly correlated to the SalePrice


	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt
Id	1.000000	0.015540	-0.014479	-0.042315	-0.058371	0.008627	-0.022610
MSSubClass	0.015540	1.000000	-0.389466	-0.197903	0.031639	-0.085553	0.021605
LotFrontage	-0.014479	-0.389466	1.000000	0.419714	0.241169	-0.047132	0.107958
LotArea	-0.042315	-0.197903	0.419714	1.000000	0.169876	-0.033113	0.028954
OverallQual	-0.058371	0.031639	0.241169	0.169876	1.000000	-0.189587	0.590761
OverallCond	0.008627	-0.085553	-0.047132	-0.033113	-0.189587	1.000000	-0.437647
YearBuilt	-0.022610	0.021605	0.107958	0.028954	0.590761	-0.437647	1.000000
YearRemodAdd	-0.030239	0.010178	0.082938	0.024308	0.568582	0.024427	0.625905
MasVnrArea	-0.072344	0.040009	0.189769	0.106600	0.419756	-0.174581	0.328897
BsmtFinSF1	-0.013234	-0.069439	0.239734	0.232341	0.230438	-0.068285	0.234207
BsmtFinSF2	0.014964	-0.073834	0.046928	0.138615	-0.081342	0.040598	-0.058987
BsmtUnfSF	-0.014316	-0.147155	0.111368	0.008924	0.297384	-0.169743	0.170077
TotalBsmtSF	-0.024541	-0.264277	0.407566	0.324476	0.547448	-0.243419	0.423763
1stFlrSF	-0.007492	-0.258207	0.453035	0.331295	0.527908	-0.166191	0.311928
2ndFlrSF	-0.005997	0.319176	0.074953	0.075311	0.265906	0.004047	-0.021327
LowQualFinSF	-0.040553	0.024935	0.010748	0.019956	-0.011186	0.047865	-0.165742
GrLivArea	-0.013772	0.078213	0.397260	0.308590	0.610102	-0.115250	0.198778

Snippet of heatmap. Complete heatmap can be found in the notebook

EDA Analysis

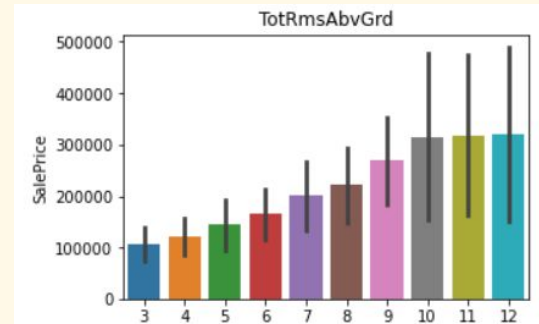
Overall Qual

High direct correlation
is observed



TotRmsAbvGrd

High direct correlation
is observed



RMSE -root mean square error

$$RMSE(\hat{\theta}) = \sqrt{MSE(\hat{\theta})} = \sqrt{E((\hat{\theta} - \theta)^2)}$$

$$RMSE = \sqrt{\frac{\sum_{i=1}^n (\hat{y}_i - y_i)^2}{n}}$$



Linear Regression

```
lreg = LinearRegression()  
lreg.fit(X_train,y_train)  
mse = mean_squared_error(y, lreg.predict(X))  
RMSE = math.sqrt(mse/num_data)  
print(RMSE)
```

939.6790554728881

Random Forest

```
RF = RandomForestRegressor(n_estimators=500,bootstrap=True,random_state=1,oob_score=True)
RF = RF.fit(X_train,y_train)
y_pred = RF.predict(X)
num_data = X.shape[0]
mse = mean_squared_error(y, y_pred)
RMSE = math.sqrt(mse/num_data)
print(RMSE)
```

403.81078260666123



Gaussian Process

```
kernel = DotProduct() + WhiteKernel()  
gp = GaussianProcessRegressor(kernel=kernel,random_state=42)  
gp.fit(X_train, y_train)  
mse = mean_squared_error(y, gp.predict(X))  
RMSE = math.sqrt(mse/num_data)  
print(RMSE)
```

1054.74356190732



Xgboost

```
xg_reg = XGBRegressor(max_depth=5, n_estimators=10)
xg_reg.fit(X_train,y_train)
mse = mean_squared_error(y, xg_reg.predict(X))
RMSE = math.sqrt(mse/num_data)
print(RMSE)
```

```
[14:19:45] WARNING: /workspace/src/objective/regression
1984.8613427728646
```

Stacking

```
def get_stacking():  
    # define the base models  
    level0 = list()  
    #level0.append(('lr', LinearRegression()))  
    level0.append(('gp', GaussianProcessRegressor(kernel=kernel,random_state=42)))  
    level0.append(('rc', RandomForestRegressor(n_estimators=500,bootstrap=True,random_state=1,oob_score=True)))  
    level0.append(('xgb', XGBRegressor(max_depth=5, n_estimators=10)))  
    # define meta learner model  
    level1 = LinearRegression()  
    # define the stacking ensemble  
    model = StackingRegressor(estimators=level0, final_estimator=level1, cv=5)  
    return model
```

```
stacked = get_stacking()  
stacked.fit(X_test, y_test)  
y_pred = stacked.predict(X)
```

```
num_data = X.shape[0]  
mse = mean_squared_error(y, y_pred)  
RMSE = math.sqrt(mse/num_data)  
print(RMSE)
```

823.7815744511447



Models used



Linear Regression

RMSE: 939.68



Gaussian Process

RMSE: 1054.74



Stacked

RMSE: 823.78



Random Forest

RMSE: 403.81



XGBoost

RMSE: 1984.86

Conclusion

Best Model
Using a Random Forest
Regressor to train the data
achieved the lowest RMSE

Stacking
The Stacking model, with
Linear Regression as its
final classifier, achieved the
next to lowest RMSE